

50.007


Machine Learning

Lu, Wei




Reinforcement Learning (I)

Bayes' Ball Algorithm




Are X and Y independent?

Bayes' Ball Algorithm




Are X and Y independent?

Bayes' Ball Algorithm




Are X and Y independent?

Markov Blanket



$$p(X|\mathbf{V}_{-X})$$


Bayesian Networks



How do we learn such probability values?

$$\begin{aligned}
 & p(x_1 = 1, x_2 = 2, x_3 = 1, x_4 = 3, x_5 = 5, x_6 = 2) \\
 &= p(x_1 = 1) \times p(x_2 = 2 | x_1 = 1) \times p(x_4 = 3 | x_1 = 1) \\
 &\times p(x_3 = 1 | x_2 = 2, x_4 = 3) \times p(x_6 = 2 | x_2 = 2, x_3 = 1) \times p(x_5 = 5 | x_4 = 3)
 \end{aligned}$$

Bayesian Networks



x_1	x_2	x_3			
		1	2	...	r_3
1	1				
1	2				
...	...				
1	r_2				
2	1				
2	2				
...	...				
2	r_2				
...	...				
r_1	1				
r_1	2				
...	...				
r_1	r_2				

What is the number of free parameters involved in this table?

(x_1, x_2)
 $(x_3 | x_1)$
 $\mathcal{C}_1)$

Overview of ML

Supervised
Learning

Unsupervised
Learning


Overview of ML

Supervised
Learning

A set of
states

Unsupervised
Learning

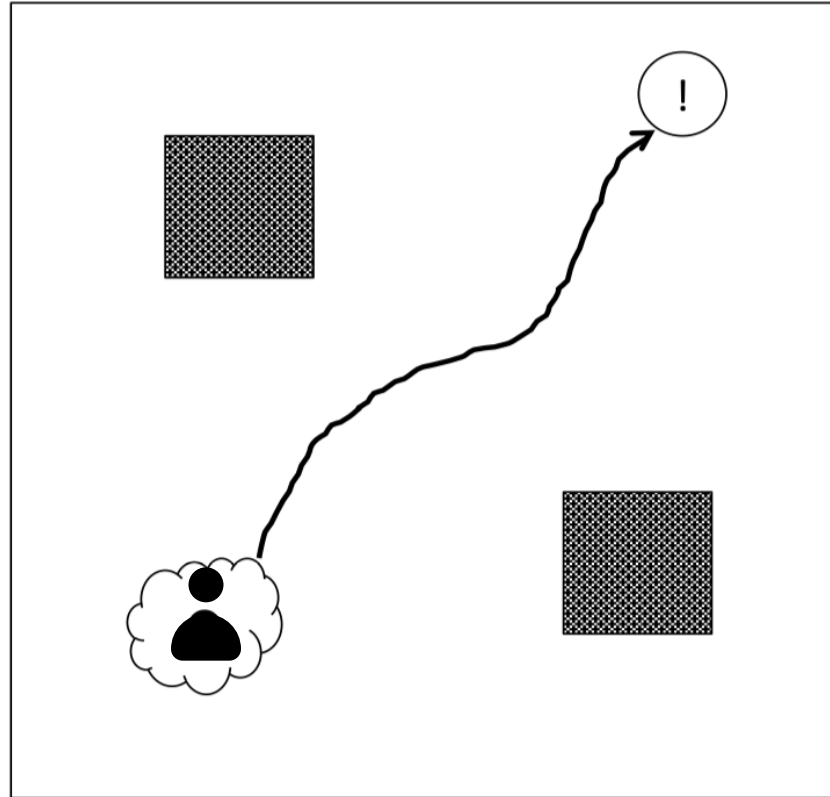
A set of
actions



$f : S \rightarrow A$

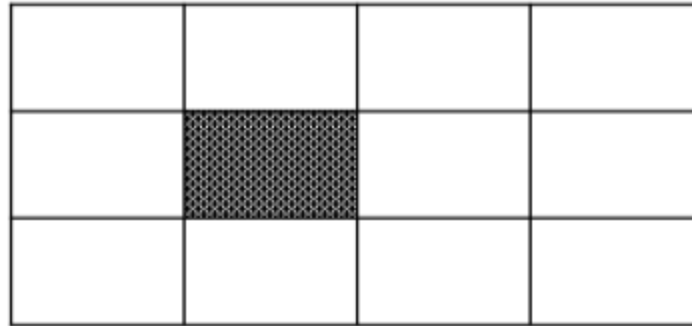
Reinforcement
Learning

Learn How to Act



How do we teach a robot how to act optimally in a complex environment?


Block World Environment



The robot can be at one block at any given time.


There is a set of states S

Block World Environment




The robot can take an action from a set of predefined possible actions at each state.

There is a set of actions A



Block World

Transition Probabilities



If the robot moves towards a particular direction, there is a 0.8 chance that it would reach the block in front, and there is a 0.1 chance to reach a state to its left (right).

A transition probability function

$$T(s, a, s') = p(s'|s, a)$$

Block World

Rewards



			+1.0
			-1.0

Each block is associated with a reward. The two blocks at the upper-right corner are assigned rewards +1 and -1.

Block World

Rewards



-0.6	+1.2	+0.1	+1.0
-0.1		-0.1	-1.0
+0.9	-0.7	-2.0	+1.3

Each block is associated with a reward. The two blocks at the upper-right corner are assigned rewards +1 and -1.

The reward is $R(s)$ for each state.

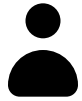
In general it can be defined as $R(s, a, s')$

action

old state

new state


Markov Decision Process



-0.6	+1.2	+0.1	+1.0
-0.1		-0.1	-1.0
+0.9	-0.7	-2.0	+1.3

- a set of states S
- a set of actions A
- a transition probability function $T(s, a, s') = p(s'|s, a)$
- a reward function $R(s, a, s')$ (or just $R(s')$)

Markov Decision Process



-0.6	+1.2	+0.1	+1.0
-0.1		-0.1	-1.0
+0.9	-0.7	-2.0	+1.3

- a set of states S
- a set of actions A
- a transition probability function $T(s, a, s') = p(s'|s, a)$
- a reward function $R(s, a, s')$ (or just $R(s')$)

How do we
learn the policy?

Block World

Utility (Long Term Reward)




-0.6	+1.2	+0.1	+1.0
0.1		-0.1	-1.0
+0.9	-0.7	-2.0	+1.3




$$U([s_0, s_1, s_2, \dots]) = R(s_0) + R(s_1) + R(s_2) + \dots = \sum_{t=0}^{\infty} R(s_t)$$

Will this be a good way of defining long term reward?

Block World

Utility (Long Term Reward)



-0.6			+1.0
0.1		-0.1	-1.0
+0.9	-0.7	-2.0	+1.3

$$U([s_0, s_1, s_2, \dots]) = R(s_0) + R(s_1) + R(s_2) + \dots = \sum_{t=0}^{\infty} R(s_t)$$

Will this be a good way of defining long term reward?

Block World

Utility (Long Term Reward)



-0.6	+1.2	+0.1	+1.0
0.1		-0.1	-1.0
+0.9	-0.7	-2.0	+1.3

$$U([s_0, s_1, s_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots = \sum_{t=0}^{\infty} \gamma^t R(s_t)$$



discount
factor

Block World

Utility (Long Term Reward)



-0.6	+1.2	+0.1	+1.0
0.1		-0.1	-1.0
+0.9	-0.7	-2.0	+1.3

$$U([s_0, s_1, s_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots = \sum_{t=0}^{\infty} \gamma^t R(s_t)$$

$$\frac{R_{min}}{1-\gamma} = \sum_{t=0}^{\infty} \gamma^t R_{min} \leq U([s_0, s_1, s_2, \dots]) \leq \sum_{t=0}^{\infty} \gamma^t R_{max} = \frac{R_{max}}{1-\gamma}$$

Lower
Bound

↓
Smallest Reward

↙
Largest Reward

Upper
Bound

Block World


Learning the Policy



-0.01	-0.01	-0.01	+1.0
-0.01		-0.01	-1.0
-0.01	-0.01	-0.01	-0.01



How to learn the policy?



→	→	→	+1.0
↓		←	-1.0
↑	→	↑	←

An example policy

Policy, Value, Q-Value

$$\pi(s)$$

a particular *policy* that specifies
the action we should take in state s .

Policy, Value, Q-Value

$$\pi(s)$$

a particular *policy* that specifies
the action we should take in state s .

$$V^\pi(s)$$

The *value* of state s under policy π .
It is the expected long-term reward of starting in state s
and act based on policy π thereafter.

Policy, Value, Q-Value

$$\pi(s)$$

a particular *policy* that specifies
the action we should take in state s .

$$V^\pi(s)$$

The *value* of state s under policy π .
It is the expected long-term reward of starting in state s
and act based on policy π thereafter.

$$Q^\pi(s, a)$$

The *Q-value* of state s and action a under policy π .
It is the expected long-term reward of starting in state s ,
taking action a and acting based on policy π thereafter.

Policy, Value, Q-Value

$$\pi^*(s)$$

The *optimal policy* $\pi^*(s)$ specifies the optimal action we should take in state s .

$$V^*(s)$$

The *value* of state s under the optimal policy π^*

$$Q^*(s, a)$$

The *Q-value* of state s and action a under the optimal policy π^*

Policy, Value, Q-Value

$$V^*(s) = ??$$

$$V^*(s)$$

The *value* of state s under the optimal policy π^*

$$Q^*(s, a)$$

The *Q-value* of state s

and action a under the optimal policy π^*

Policy, Value, Q-Value

$$V^*(s) = \max_a Q^*(s, a)$$

$$V^*(s)$$

The *value* of state s under the optimal policy π^*

$$Q^*(s, a)$$

The *Q-value* of state s

and action a under the optimal policy π^*

Policy, Value, Q-Value

$$V^*(s) = \max_a Q^*(s, a) = Q^*(s, \pi^*(s))$$

$$\pi^*(s)$$

The *optimal policy* $\pi^*(s)$ specifies the optimal action we should take in state s .

$$V^*(s)$$

The *value* of state s under the optimal policy π^*

$$Q^*(s, a)$$

The *Q-value* of state s and action a under the optimal policy π^*

Policy, Value, Q-Value

$$Q^*(s, a) = ??$$

$$\pi^*(s)$$

The *optimal policy* $\pi^*(s)$ specifies the optimal action we should take in state s .

$$V^*(s)$$

The *value* of state s under the optimal policy π^*

$$Q^*(s, a)$$

The *Q-value* of state s and action a under the optimal policy π^*

Policy, Value, Q-Value

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

probability of next state	immediate reward	long-term reward
------------------------------	---------------------	---------------------

$$\pi^*(s)$$

The *optimal policy* $\pi^*(s)$ specifies
the optimal action we should take in state s .

$$V^*(s)$$

The *value* of state s under the optimal policy π^*

$$Q^*(s, a)$$

The *Q-value* of state s
and action a under the optimal policy π^*

Policy, Value, Q-Value

$$V^*(s) = \max_a Q^*(s, a)$$


$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$\pi^*(s)$$

The *optimal policy* $\pi^*(s)$ specifies the optimal action we should take in state s .

$$V^*(s)$$

The *value* of state s under the optimal policy π^*

$$Q^*(s, a)$$

The *Q-value* of state s and action a under the optimal policy π^*

Policy, Value, Q-Value

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$\pi^*(s)$$

The *optimal policy* $\pi^*(s)$ specifies the optimal action we should take in state s .

$$V^*(s)$$

The *value* of state s under the optimal policy π^*

$$Q^*(s, a)$$

The *Q-value* of state s and action a under the optimal policy π^*

Policy, Value, Q-Value

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

Updating Value

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

Value Iteration

1. Start with $V_0^*(s) = 0$, for all $s \in S$
2. Given V_i^* , calculate the values for all states $s \in S$

$$V_{i+1}^*(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i^*(s')]$$

3. Repeat the above until convergence

Value Iteration

1. Start with $V_0^*(s) = 0$, for all $s \in S$
2. Given V_i^* , calculate the values for all states $s \in S$

$$V_{i+1}^*(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i^*(s')]$$

3. Repeat the above until convergence

There is a guarantee that this process will converge

Question

How do we recover the optimal policy based on the values?

Policy, Value, Q-Value

$$V^*(s) = \max_a Q^*(s, a) = Q^*(s, \pi^*(s))$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

Policy, Value, Q-Value

$$V^*(s) = \max_a Q^*(s, a) = Q^*(s, \pi^*(s))$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

Policy, Value, Q-Value

$$V^*(s) = \max_a Q^*(s, a) = Q^*(s, \pi^*(s))$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

Policy, Value, Q-Value

$$\pi^*(s) = ??$$

$$\pi^*(s)$$

The *optimal policy* $\pi^*(s)$ specifies the optimal action we should take in state s .

$$Q^*(s, a)$$

The *Q-value* of state s and action a under the optimal policy π^*

Policy, Value, Q-Value

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$

$$\pi^*(s)$$

The *optimal policy* $\pi^*(s)$ specifies
the optimal action we should take in state s .


$$Q^*(s, a)$$

The *Q-value* of state s
and action a under the optimal policy π^*

Learning Optimal Policy

Step 1


Run Value Iteration Algorithm



Step 2

Calculate the Q values

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$



Step 3

Find the optimal action for each state


$$\pi^*(s) = \arg \max_a Q^*(s, a)$$

Value Iteration

$$R(s) = -0.01$$



-0.01	-0.01	-0.01	+1.0
-0.01		-0.01	-1.0
-0.01	-0.01	-0.01	-0.01




			+1.0
			-1.0

Value Iteration

Learned Policy



-0.01	-0.01	-0.01	+1.0
-0.01		-0.01	-1.0
-0.01	-0.01	-0.01	-0.01




→	→	→	+1.0
↑		←	-1.0
↑	←	←	↓


Value Iteration

$$R(s) = -2.0$$



-2.0	-2.0	-2.0	+1.0
-2.0		-2.0	-1.0
-2.0	-2.0	-2.0	-2.0




			+1.0
			-1.0

Value Iteration

Learned Policy



-2.0	-2.0	-2.0	+1.0
-2.0		-2.0	-1.0
-2.0	-2.0	-2.0	-2.0




→	→	→	+1.0
↑		→	-1.0
→	→	→	↑

Value Iteration

Learned Policy



-2.0	-2.0	-2.0	+1.0
-2.0		-2.0	-1.0
-2.0	-2.0	-2.0	-2.0



Why?



→	→	→	+1.0
↑		→	-1.0
→	→	→	↑

Value Iteration

Step 1

Since the procedure relies on Q -values,
is it possible to design an algorithm
that directly computes these Q -values?

Step 2

Find the optimal action for each state

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$

Value Iteration

Step 1

Since the procedure relies on Q -values,
is it possible to design an algorithm
that directly computes these Q -values?

We will discuss Q -value
iteration next time!

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$